Università degli Studi di Verona

Department of Computer Science

Internal Internship Project Report

Skill Gap Evaluation and Course Recommendations

Author:
Simone Pedretti
Student ID: VR500040

Academic Year 2025/2026

Contents

1 Introduction **2**

2 Individual User Features **3**
2.1 Home . . . 3
2.1.1 Browsing for Roles and Skills . . . 3
2.1.2 Project Manager . . . 3
2.2 Profile . . . 4
2.2.1 Target Roles and Skills Management . . . 4
2.2.2 Cedefop Forecast and Skill Gap Evaluation . . . 5
2.2.3 Course Recommendation . . . 5

3 Organization Features **6**
3.1 Home . . . 6
3.1.1 Projects . . . 6
3.2 Profile . . . 6
3.2.1 Team Members . . . 6
3.2.2 Global Skill Gap Evaluation . . . 6
3.2.3 Courses Management . . . 6

4 Guest Features **8**

5 ESCO API Integration **9**

6 Data Management **9**
6.1 Database Schema . . . 9
6.2 Main Queries and Data Processing Logic . . . 9

7 Testing and Quality Assurance **11**

1 Introduction

This report details the work carried out during my internal internship at *Università degli Studi di Verona*. The primary objective of this project was the design and development of a web-based platform tailored for skill management, competence tracking, and course recommendation.

From a technical perspective, the application follows a standard client-server architecture, with a clear separation of concerns between the frontend and the backend. The backend system, hosted on the server to handle the core business logic and data management, was developed in Python utilizing the **FastAPI** framework. It heavily relies on the **JSON** format for efficient data structuring and API communication. The frontend system, on the client side, was implemented using standard web technologies, specifically **HTML** and **CSS**. This ensures a clean, responsive, and accessible experience for both individual users and organizational profiles interacting with the system.

The following sections will provide a detailed overview of the specific features implemented for the different types of users, detailing the browsing capabilities, profile management, and the skill gap evaluation mechanisms. Furthermore, the report will explore the underlying technical components that power these functionalities: final chapters will discuss the integration of the **ESCO API** for standardized competence mapping, the strategies adopted for data management, and finally, the testing procedures conducted to ensure the reliability and robustness of the system.

2 Individual User Features

This section outlines the functionalities available to individual users interacting with the platform. Within this system, a user can hold one of two distinct roles: a standard Employee or a Manager. While an Employee utilizes the platform primarily for personal skill tracking and career development, a Manager is a user who has been elevated to this status by their affiliated organization. Managers inherit all standard employee capabilities but are additionally granted administrative privileges to oversee projects and coordinate team members within their company.

2.1 Home

The Home dashboard serves as the primary entry point for users, offering intuitive search and management tools depending on their assigned role.

2.1.1 Browsing for Roles and Skills

Users can explore various professional profiles through the role browsing feature. This search mechanism integrates the ESCO API, ensuring that role definitions, responsibilities, and terminologies are standardized and perfectly aligned with the European labor market classifications.

Similarly to the role search, users can browse for individual competencies. By querying the ESCO API, the system retrieves a list of macthing skills. This allows users to accurately identify specific abilities.

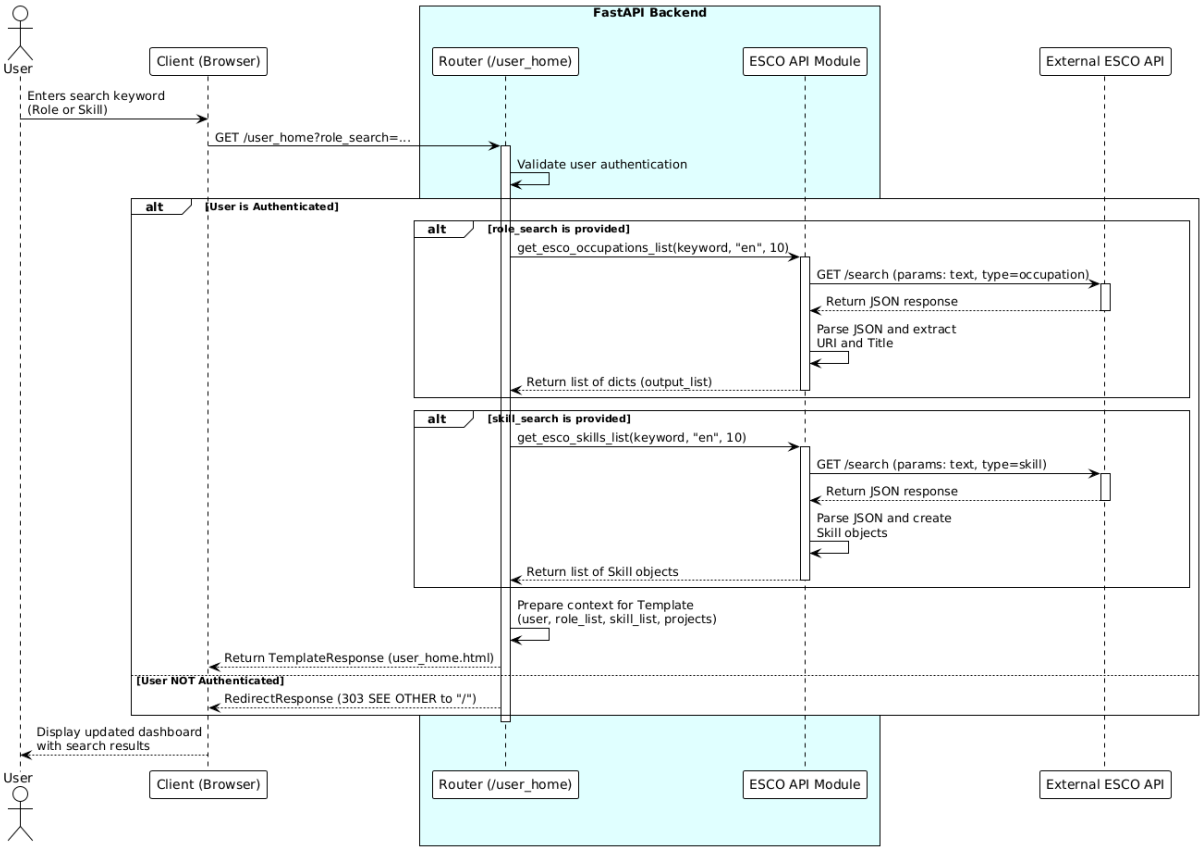

Figure 1: Sequence Diagram: User - Esco API

2.1.2 Project Manager

This dedicated area is exclusively accessible to users with Manager privileges. Here, managers can create, organize, and oversee various internal projects. They have the authority to assign team members to these projects, provided those individuals belong to the same organization. Furthermore, managers can select specific target roles for a project to calculate Cedefop-based employment forecasts and generate

preliminary course recommendations for the team. A more comprehensive explanation of the forecasting and recommendation mechanisms is detailed in the Profile section below.

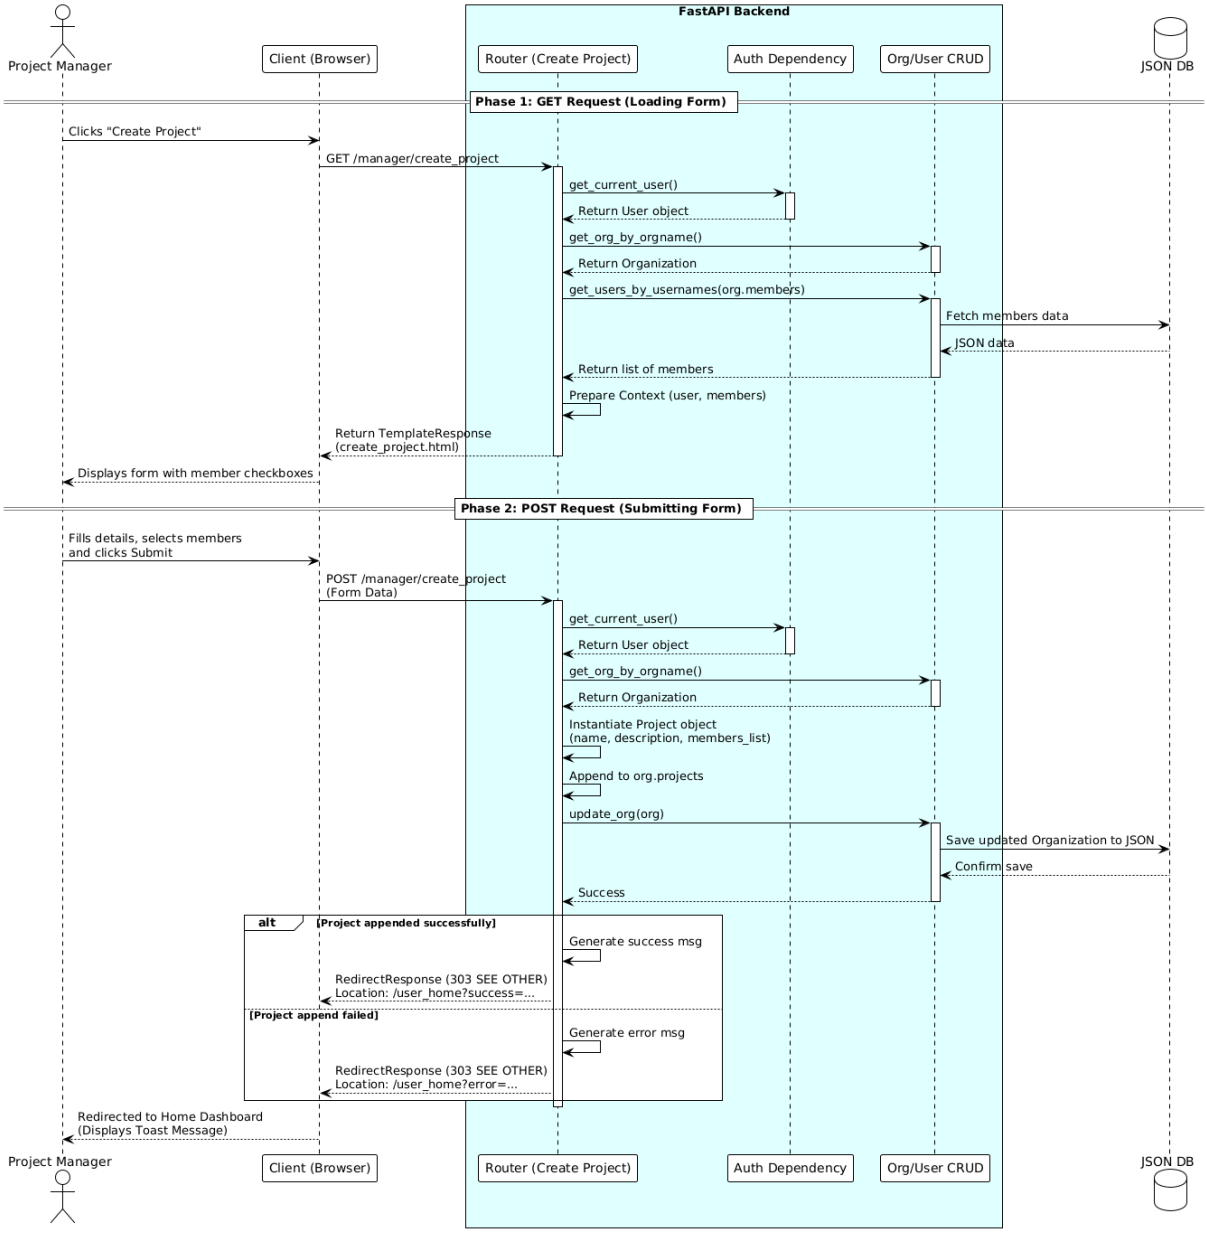

Figure 2: Sequence Diagram: Project Creation

2.2 Profile

The Profile section acts as the personal hub for the user, centralizing account settings, competence tracking, and career development tools.

2.2.1 Target Roles and Skills Management

Both employees and managers can access this area to view their personal data and securely change their passwords. More importantly, users can manage their professional portfolio by viewing, adding, or removing their personal skills and desired target roles. To streamline the data entry process, the platform supports bulk operations: users can download a standardized CSV template directly from the profile interface, enabling the simultaneous upload of multiple skills. Furthermore, the dashboard provides a

direct link to the external *MuchSkills* platform, serving as an auxiliary tool to help users intuitively map and build their personal competence profile before importing it into the system.

If a user is currently affiliated with an organization, a dedicated subsection displays their "organizational skills"—competencies that are directly assigned and managed by the company. Through this interface, users also retain the autonomy to officially leave their current organization if needed.

2.2.2 Cedefop Forecast and Skill Gap Evaluation

Once a user has defined their current skill set and selected one or more target roles, the platform can perform a comprehensive skill gap analysis. The system compares the user's existing competencies against the specific requirements of the desired roles, highlighting areas for improvement. Additionally, the application retrieves Cedefop forecast data related to these target roles, offering the user valuable insights into future labor market trends, job demand, and employment projections.

2.2.3 Course Recommendation

Based on the exact skill gaps identified in the previous step, the platform provides customized educational suggestions. This recommendation engine—which is the same core logic utilized in the Project Manager dashboard—analyzes the missing competencies and cross-references them with available training resources. By doing so, the system generates a highly personalized learning path, proposing the most relevant courses to help the user efficiently bridge the gap between their current capabilities and their professional goals.

To facilitate offline review and long-term career planning, users can conveniently export the complete analysis—including the skill gap results and course recommendations—into a consolidated PDF report.

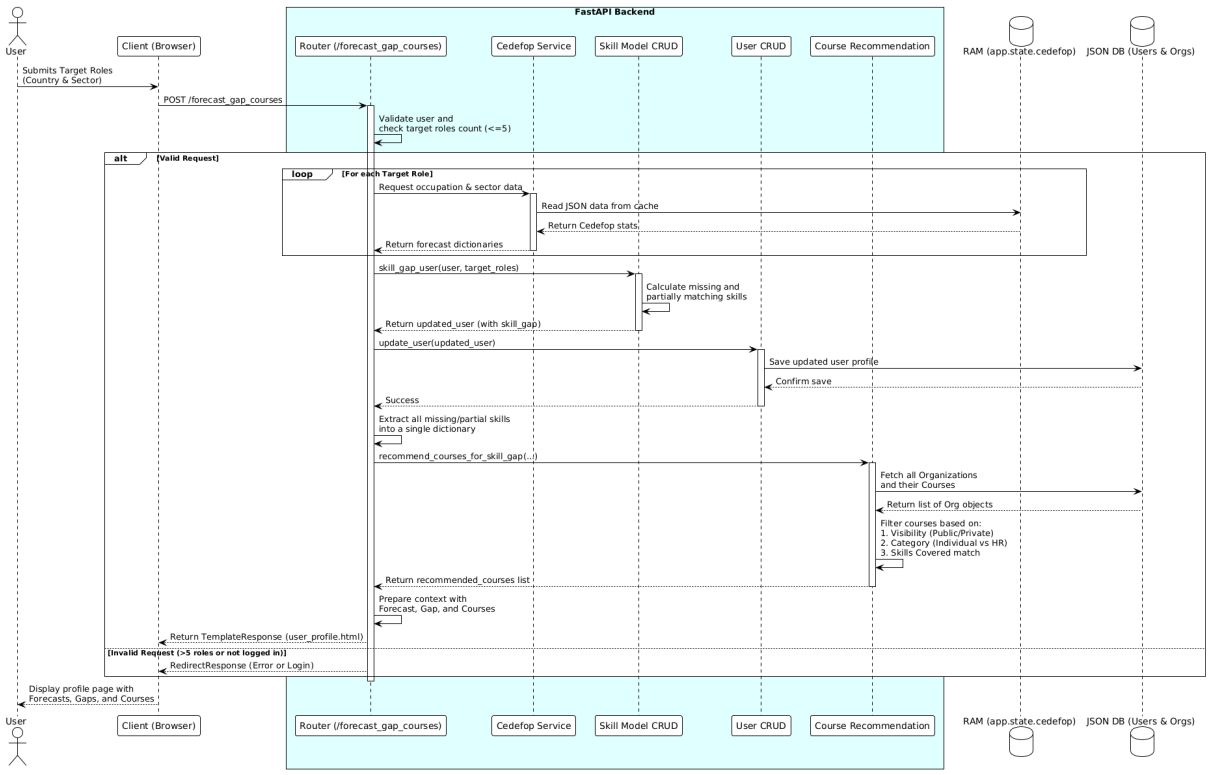

Figure 3: Sequence Diagram: Cedefop's Forecast, Skill Gap and Course Recommendation

3 Organization Features

This section details the functionalities available to Organization accounts. Unlike individual users who focus on personal career development, an Organization profile represents the company as a whole. Its primary purpose is to oversee the workforce, manage internal training resources, and analyze competence gaps at a macro level, providing an administrative overview of the company's human resources.

3.1 Home

The Home dashboard for an organization is designed to provide a high-level summary of ongoing activities, focusing primarily on project monitoring.

3.1.1 Projects

Through this interface, the organization can access a comprehensive list of all active projects currently overseen by its internal Managers. It is important to note that this is a strictly read-only view: the organization can monitor the status, but the direct management, creation, and modification of the projects remain the sole responsibility of the delegated Managers.

3.2 Profile

The Profile area serves as the central administrative hub for the organization, allowing the management of account credentials, personnel, and proprietary training materials. As with individual users, this section also includes basic settings to view corporate data and securely update the account password.

3.2.1 Team Members

The organization can view its current workforce, distinguishing between standard employees and managers. To expand the team, the organization can send an official invitation by entering a specific username into the recruitment form, facilitated by a comprehensive list of all registered platform users. The invited user then retains the autonomy to accept or decline the request. Furthermore, the organization holds the authority to adjust user privileges within the company, allowing them to promote standard employees to the Manager role or demote them as organizational needs dictate. In addition, the platform provides a direct link to the external *MuchSkills* platform. This allows the organization to design a comprehensive skill model for its workforce and seamlessly integrate these profiles into the system via a bulk CSV import.

3.2.2 Global Skill Gap Evaluation

Provided the organization has active projects, it can perform a macroscopic analysis of its workforce capabilities. Instead of evaluating a single user, this feature calculates a global skill gap by comparing the aggregated competencies of the current employees against the collective requirements of the ongoing projects. Based on these company-wide deficiencies, the platform suggests HR-oriented courses. These specific recommendations are designed for strategic corporate upskilling and human resource planning, rather than individual learning paths. Similar to the functionality available to individual users, organizations can export this comprehensive evaluation into a consolidated PDF report, facilitating offline review and strategic decision-making at the management level.

3.2.3 Courses Management

The platform allows organizations to act as educational providers by creating and distributing their own training modules. Through this interface, the company can view a complete catalog of its authored courses, seamlessly create new ones, and edit or delete existing materials. A critical feature of this management tool is the visibility toggle: courses can be published as "public" to be accessible by any user on the platform, or set to "private," ensuring that only the organization's officially affiliated employees can view and benefit from them.

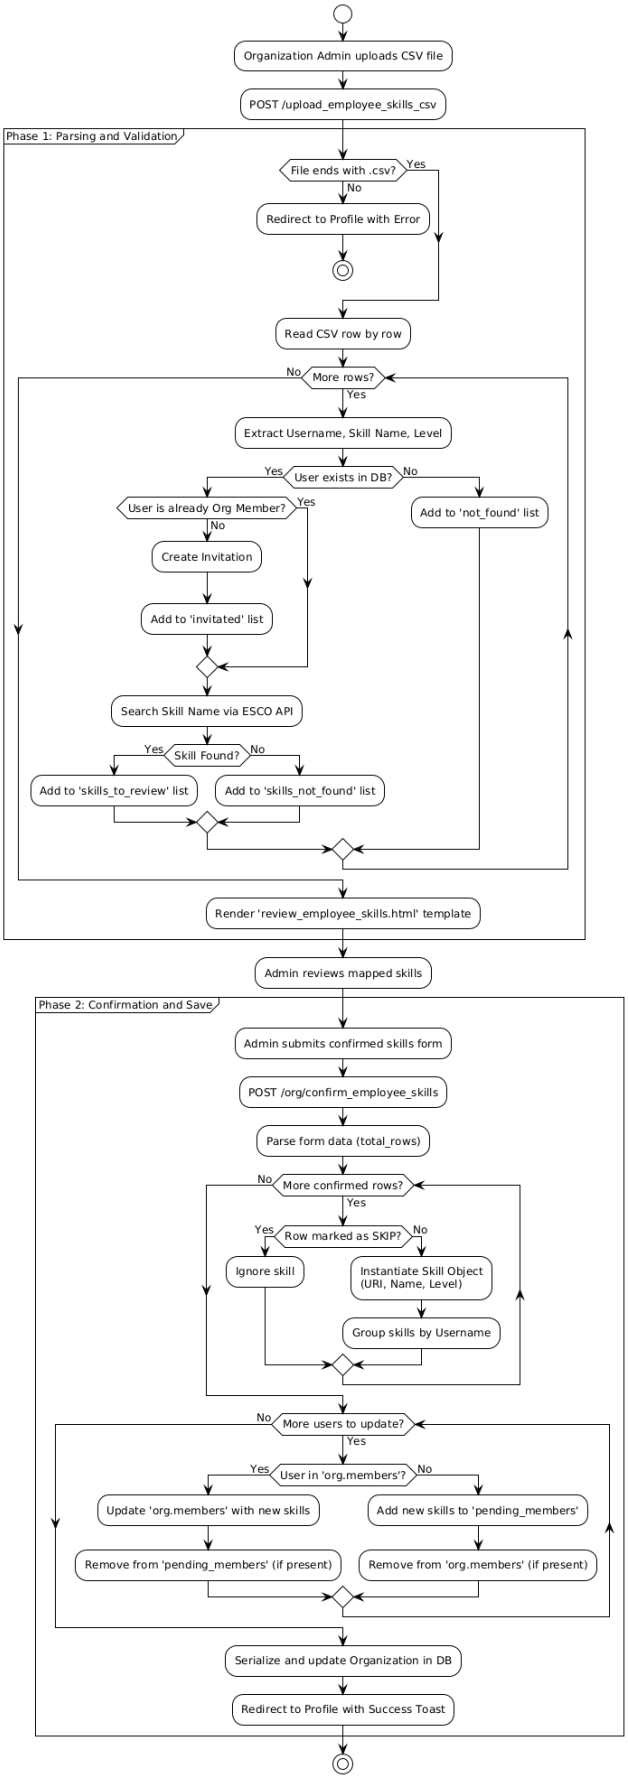

Figure 4: Activity Diagram: CSV Upload and Skill confirmation

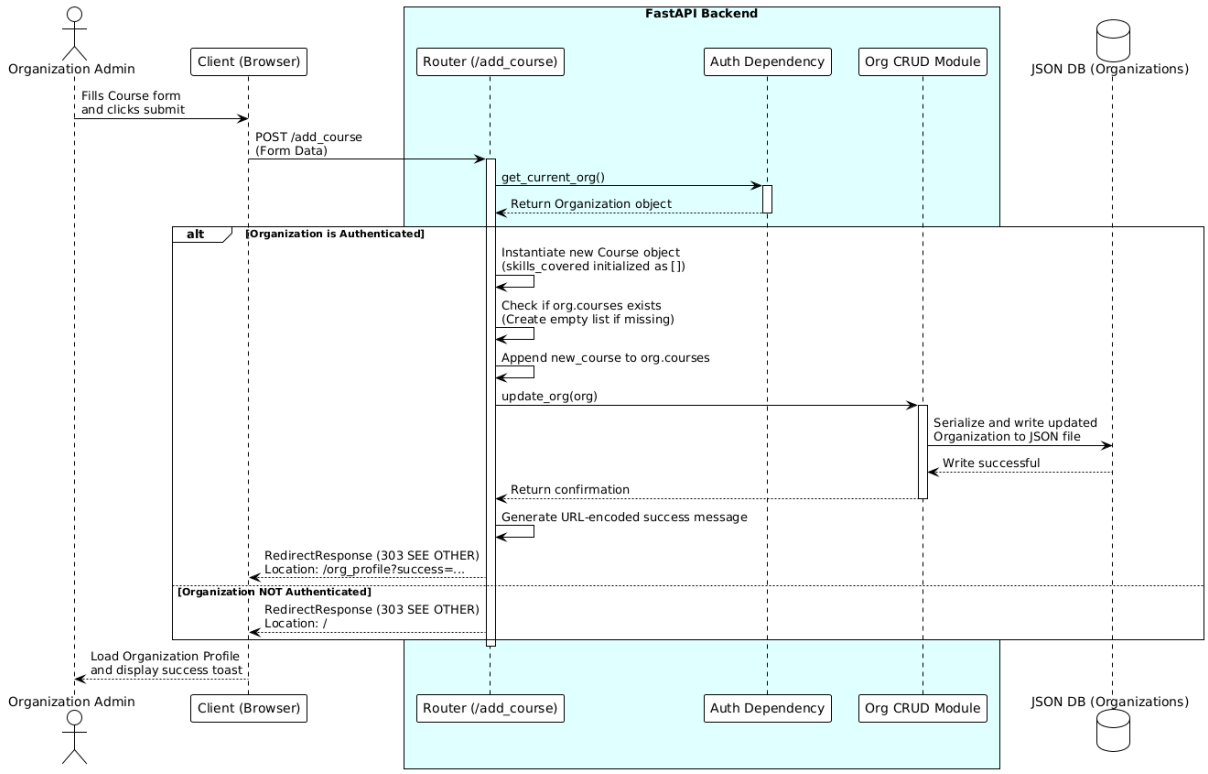

Figure 5: Sequence Diagram: Course Creation

4 Guest Features

The platform also accommodates unregistered users by providing a restricted "Guest" access mode. For individuals who prefer not to create a formal account, the system offers a streamlined interface to explore the core competency catalog. Through this guest portal, users can perform basic searches for professional roles and individual skills, which are directly powered by the ESCO API.

However, compared to registered employees or organizations, this browsing experience is intentionally synthesized and limited in scope. Guests can view roles of the queried terms, but they cannot access advanced, stateful features such as profile creation, skill gap evaluation, or personalized course recommendations. This approach allows external users to evaluate the platform's core search capabilities while maintaining data privacy and encouraging formal registration for full access.

5 ESCO API Integration

A core component of the platform is the integration with the ESCO API, hosted by the European Commission. The system interacts with this external service using the Python `requests` library to fetch standardized data regarding professional roles and specific skills.

The backend implements dedicated functions to query the API for single role details (including the ISCO classification code and essential skills), as well as search functions to retrieve lists of occupations or skills based on user input. For instance, the system requests the `hasEssentialSkill` relation from the API to automatically map the necessary competencies for a given job title. The JSON responses provided by the API are subsequently parsed and converted into internal `Role` and `Skill` data models to be utilized by the application's business logic.

6 Data Management

Given the platform's reliance on extensive datasets, particularly for the Cedefop employment forecasts, efficient data management was a priority. The original Cedefop data, initially provided in Excel format, was converted into structured JSON files.

To ensure high performance and minimize read operations from the disk during execution, the application utilizes the FastAPI `app.state` object. Upon server startup, the system loads these JSON files directly into the server's RAM. This in-memory caching strategy significantly reduces the latency when users calculate skill gaps and request forecast data, allowing the application to respond to these computationally heavy requests almost instantaneously. User and organization data, as well as active invitations, are also managed via JSON-based local databases, allowing for a lightweight and easily deployable architecture.

6.1 Database Schema

While the platform does not use a traditional relational database, its schema is explicitly defined and validated using Python `Pydantic` models. This approach ensures strict typing, data validation, and rapid serialization/deserialization when interacting with the JSON files. The complete database schema is centrally defined within the `app/models.py` file of the project's source code. The primary entities structuring the platform are:

- **Skill and Role:** These are the foundational building blocks of the platform, designed to perfectly mirror the data structure retrieved from the external ESCO API.
- **User Model:** Stores individual credentials, assigned roles (defined by the `UserLevel` Enum), and organizational affiliation. It relies heavily on nested data, storing the user's personal portfolio in the `individual_skills` array, alongside the selected `target_roles` and the dynamically calculated `skill_gap` objects.
- **Organization Model:** Represents corporate entities. It contains the organization's credentials and dynamically embeds comprehensive lists of authored `courses` and internal `projects`. Workforce mapping is ingeniously handled via dictionaries (`members` and `pending_members`), mapping string usernames directly to their respective arrays of `Skill` objects.
- **Project and Course Models:** Embedded within the Organization model, the `Project` tracks the `manager`, `assigned_members`, and specific `target_roles`. The `Course` model tracks detailed metadata (e.g., `ects`, `category`, `cost`) and the `skills_covered` required for the recommendation engine. The generation of unique identifiers is automated using the `uuid4` standard.
- **Invitation Model:** An independent structure that manages the asynchronous flow of adding users to organizations, tracking the `orgname`, `username`, and the `status` (pending, accepted, declined) of the request.

6.2 Main Queries and Data Processing Logic

Given this file-based and in-memory architecture, data retrieval does not rely on standard SQL statements. Instead, "queries" are implemented as optimized Python CRUD (Create, Read, Update, Delete) operations that interact with the in-memory data structures or JSON files. The principal query mechanisms and algorithmic logic utilized by the system include:

- **Direct File I/O (CRUD Operations):** Modules like `crud_user.py` and `crud_org.py` manage persistent storage by mapping objects to physical JSON files. Unique identifiers, such as the `username` or `orgname`, act as primary keys to construct file paths (e.g., `data/users/{username}.json`). Updates are performed securely by serializing the Pydantic models back into the file system using `model_dump_json()`.
- **Relational Mapping and Aggregation:** To emulate SQL `JOIN` operations, the system uses aggregation functions. For instance, `get_users_by_usernames` iteratively queries the file system to resolve an array of string usernames into a fully populated list of `User` objects. This is vital for the Project Manager dashboard to visualize team members.
- **Skill Gap Algorithmic Queries:** The `crud_skill_models.py` module executes the most computationally intensive "queries". To optimize performance, the algorithm first flattens the user's current competencies into a hash map (`user_skills_dict`), enabling O(1) time complexity lookups via the skill URI. The mathematical engine then evaluates the gap: a fully possessed skill grants 1.0 point, whereas partially matched skills calculate a fractional score based on the ratio of the user's level versus the required level (`user_level / req_level`). The output categorizes skills into `matching`, `partially_matching`, and `missing` arrays.
- **In-Memory Navigation (Cedefop Data):** The `cedefop_read.py` module handles data extraction from the Cedefop JSON loaded entirely into the server's RAM during startup. Instead of disk reads, it utilizes nested dictionary `.get()` methods chained together. The module features smart parsing logic: it dynamically slices the incoming ISCO code string to navigate the correct hierarchical depth of the dataset, such as extracting the first two digits (`isco_clean[:2]`) to query sector employment details or qualifications.

7 Testing and Quality Assurance

To guarantee the stability and reliability of both the backend calculations and the frontend endpoints, a comprehensive testing suite was implemented using the **pytest** framework. The testing strategy covers HTML route rendering, internal backend logic (such as the skill gap evaluation algorithms), and proper HTTP status code responses.

To simulate client interactions without requiring a live server, the FastAPI `TestClient` is utilized. Furthermore, to isolate the testing environment and prevent the corruption of actual production data, the system employs the `monkeypatch` utility provided by pytest. This allows the automatic generation of temporary, isolated JSON directories for users, organizations, and invitations during the testing phase, which are automatically discarded upon completion.

The following snippet illustrates the configuration of the unified mocking database used across the test suite:

```python
import pytest
from fastapi.testclient import TestClient
from app.main import app
from app.crud import crud_user, crud_org

@pytest.fixture(scope="module")
def client():
    with TestClient(app) as c:
        yield c

@pytest.fixture(autouse=True)
def mock_database(tmp_path, monkeypatch):
    # Create temporary directories
    temp_users_dir = tmp_path / "test_users"
    temp_orgs_dir = tmp_path / "test_orgs"
    temp_inv_dir = tmp_path / "test_invitations"

    temp_users_dir.mkdir(exist_ok=True)
    temp_orgs_dir.mkdir(exist_ok=True)
    temp_inv_dir.mkdir(exist_ok=True)

    # Forcing code to use new tmp dirs
    monkeypatch.setattr(crud_user, "DATA_DIR_USERS", str(temp_users_dir))
    monkeypatch.setattr(crud_user, "DATA_INV_DIR", str(temp_inv_dir))

    monkeypatch.setattr(crud_org, "DATA_DIR_ORGS", str(temp_orgs_dir))
    monkeypatch.setattr(crud_org, "DATA_INV_DIR", str(temp_inv_dir))

    yield
```

Listing 1: Pytest configuration with TestClient and temporary mock databases